Sudhindra Kumar

Email: kumarsudhindra2910@gmail.com
LinkedIn: linkedin.com/sudhindra-kumar
Mobile: +91-9546602057
GitHub: github.com/sudhindra62

Professional Summary

Dynamic AI/ML Engineer and Full-Stack Developer specializing in Agentic AI, Data Structures & Algorithms, and Retrieval-Augmented Generation (RAG). Expertise in designing and deploying scalable software solutions, utilizing REST APIs, System Design, CI/CD practices, and Cloud Computing to drive efficiency and performance. Three-time hackathon winner recognized for the ability to architect production-ready applications across AI, mobile, gaming, and full-stack domains, with a focus on optimization, automation, monetization, and seamless end-to-end delivery. Committed to driving innovation and enhancing user experiences through cutting-edge technology solutions that push the boundaries of what is possible.

Education

Maharaja Institute of Technology Mysore (Mysore, Karnataka) — Aug 2023 – Jul 2027
B.Tech — Artificial Intelligence & Machine Learning — *CGPA: 8.8/10*
Cleared JEE Main and secured admission offers from multiple NITs — 2023

Skills Summary

- **Software Engineering:** Java Expertise, R, Python, C, C #, Node.js Development, Android Software Engineering, Firebase Integration, Containerization (Docker) expertise, Experience with Cloud Platforms, Data Structure Analysis
- **AI/ML:** Machine Learning, Sentiment Analysis, Neural Network Design, Agentic AI, RAG, LLM Fine-Tuning, Predictive Analytics, TensorFlow, Langchain, LangGraph, Scikit-learn, Image Processing using OpenCV
- **Databases & Data Engineering:** Data Manipulation with SQL, PostgreSQL, NoSQL,

Experience

Full-Stack Software Engineer — Navodita Infotech — Jan 2026 – Jun 2026
React.js, Node.js, Python, REST APIs, Docker, CI/CD, Cloud Deployment — *Full-Stack & Backend Engineering*

- **Deployment Efficiency**: Developed and implemented scalable applications while streamlining workflows, resulting in a **35% increase in deployment efficiency** through improved CI/CD execution and modular application design.
- **Application Throughput**: Executed API optimization and CI/CD automation initiatives, resulting in a **28% improvement in application throughput** and stronger production performance across backend services.

Backend Software Engineer — JP Morgan Chase & Co. — Mar 2026 – May 2026
Apache Kafka, H2 Database, RESTful APIs, Microservices, Java, Distributed Systems *Backend Engineering & Distributed Systems*

- **Event-Driven Architecture & Testing**: Architected high-throughput RESTful backend services with Apache Kafka for distributed event streaming, attaining a **76% improvement in system reliability**; led unit, integration, and microservice communication validation using H2 in-memory database, ensuring zero-defect regression across production releases.

Achievements

Three-Time Hackathon Winner
Competitive Programming, Full-Stack Delivery, Agentic AI

- **Game Dev**: Secured 1st place among 152 teams by optimizing architecture and improving application responsiveness by 30%.
- **Full-Stack Win**: Won Full-Stack Hackathon among 42+ teams by delivering scalable Android and backend solutions.
- **Agentic AI**: Achieved 84% predictive accuracy with 47-second inference performance in Agentic AI competition.

Products

ClarityBoost — AI Media Enhancement App — Kotlin, Compose, Firebase, OpenCV, FFmpeg
Android AI, Offline Processing — *Billing • Analytics • AdMob*

- **Offline AI Enhancement**: Enabled offline AI media enhancement with GPU-accelerated processing pipelines.
- **Quality Improvement**: Improved media enhancement quality by 60%+ through OpenCV and FFmpeg optimization.
- **Monetization Readiness**: Increased deployment readiness and monetization capability through Billing, Analytics, AdMob, and workflow automation.

Projects

Disaster Response Nexus [GitHub] — React.js, TypeScript, Agentic AI
AI-Driven Disaster Management — *Predictive Analytics, Multi-Agent Orchestration*

- **Prediction Accuracy**: Achieved 84% prediction accuracy by engineering an AI-driven disaster response platform using predictive analytics and multi-agent orchestration.
- **Response Time**: Reduced response generation time upto 47 seconds through optimized autonomous decision pipelines and AI inference workflows.

Endless Runner Mobile Game [GitHub] — Unity, C#, Firebase, Android
Mobile Game Development — *SDLC, Profiling, Memory Optimization*

- **Performance**: Delivered a production-ready game sustaining 60 FPS through modular architecture and memory optimization.
- **Runtime Efficiency**: Improved runtime efficiency and gameplay responsiveness through SDLC-driven optimization and performance engineering.

Certifications

Applied Machine Learning in Python — Introduction to Cloud Computing — Mathematics for Machine Learning — Statistics for Data Science — Data Science Fundamentals with Python and SQL — Blockchain and Cryptography Overview
Python for AI & Data Science — Generative AI Applications — Create Machine Learning Models with Microsoft Azure — Introduction to DevOps — AI Agents builder — Agentic AI and AI Agents :A Primer for Leaders
ML Pipelines with Azure ML Studio — Linear Algebra for ML and AI — JP Morgan Software Engineering Simulation